

Logistic Regression

16.07.2014

Markus Holzemer
Jonas Traub
Timo Walther

Agenda

Motivation

Logistic Regression Algorithm

Parallelization Strategy

Implementation in Stratosphere and Spark

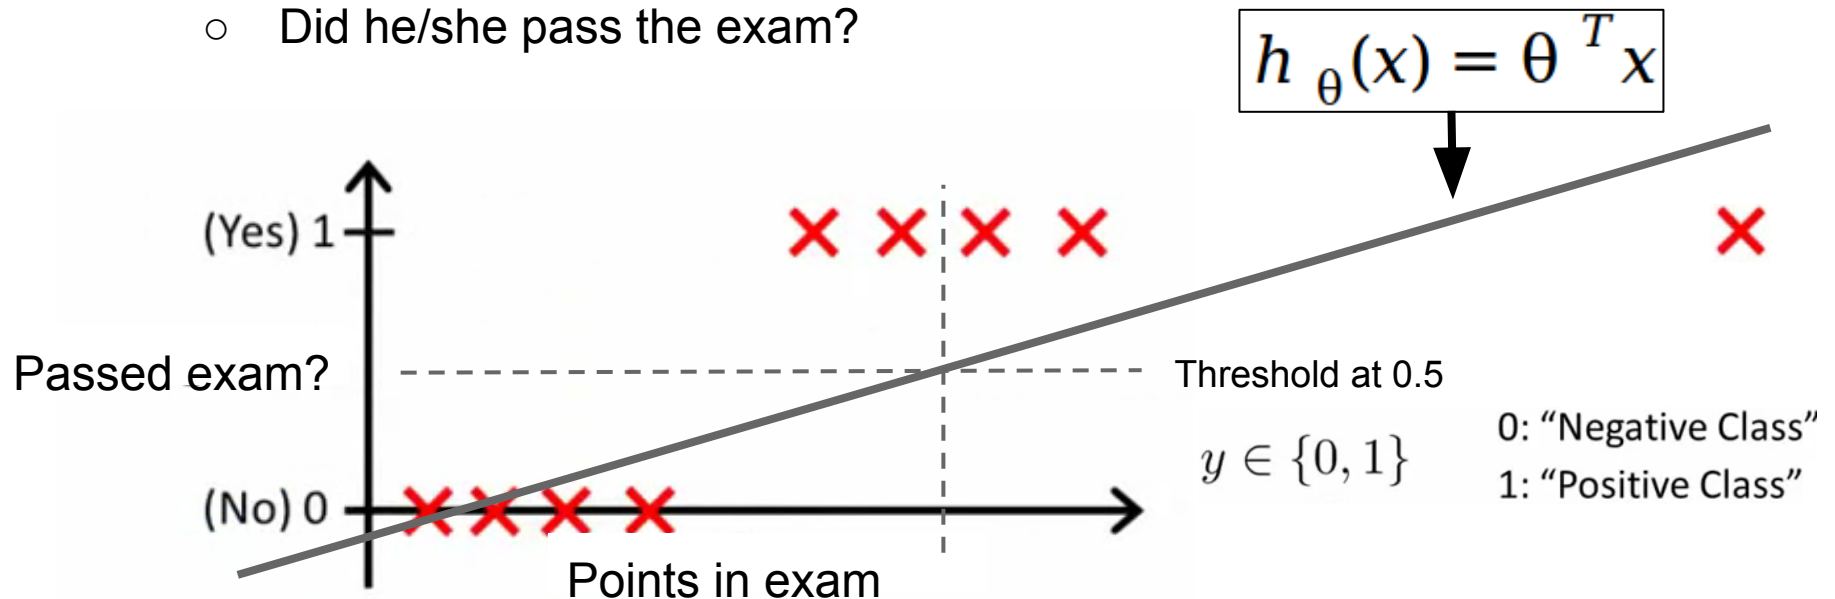
Experiment Results

Motivation

- Logistic Regression is for **Classification**
- Typically binary classification
 - Is this mail spam?
 - Did he/she pass the exam?

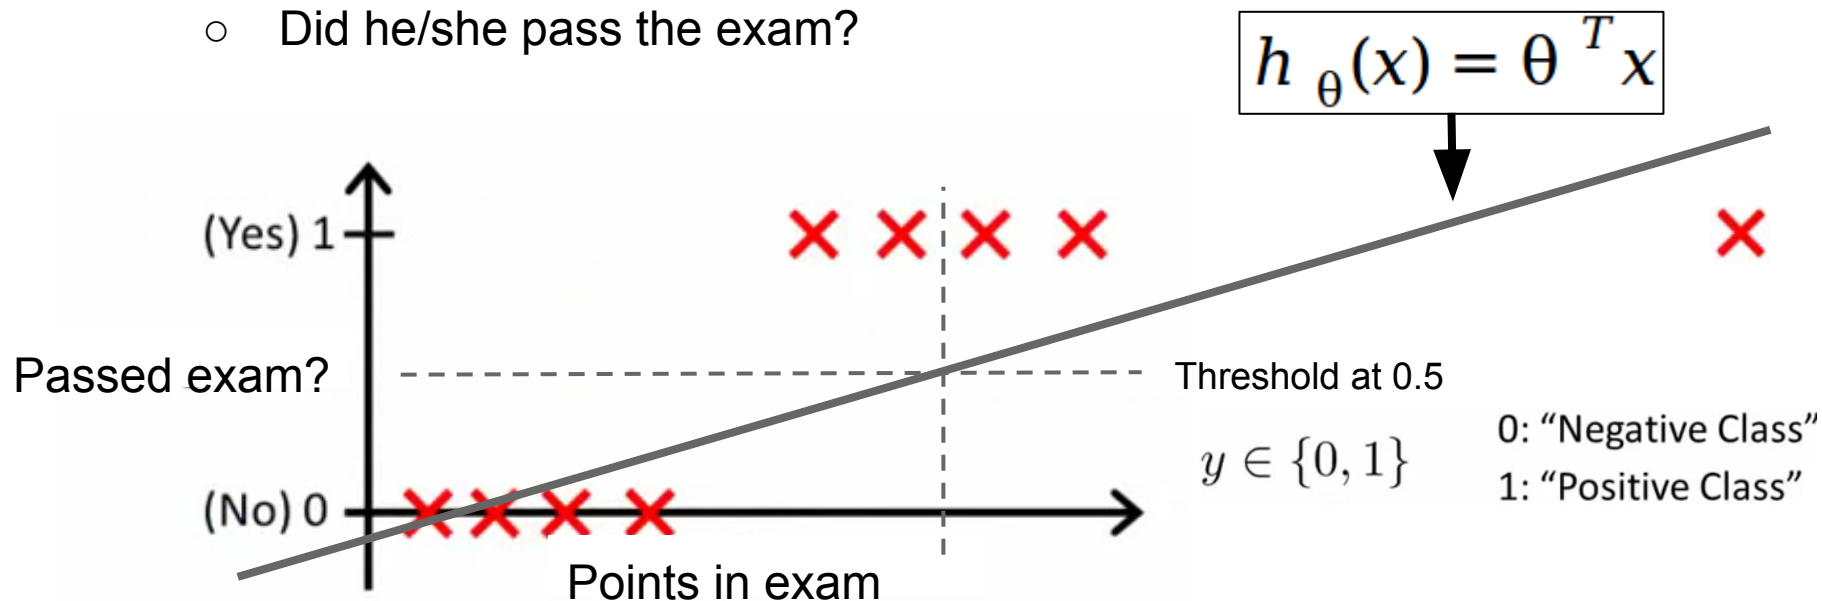
Motivation

- Logistic Regression is for **Classification**
- Typically binary classification
 - Is this mail spam?
 - Did he/she pass the exam?



Motivation

- Logistic Regression is for **Classification**
- Typically binary classification
 - Is this mail spam?
 - Did he/she pass the exam?



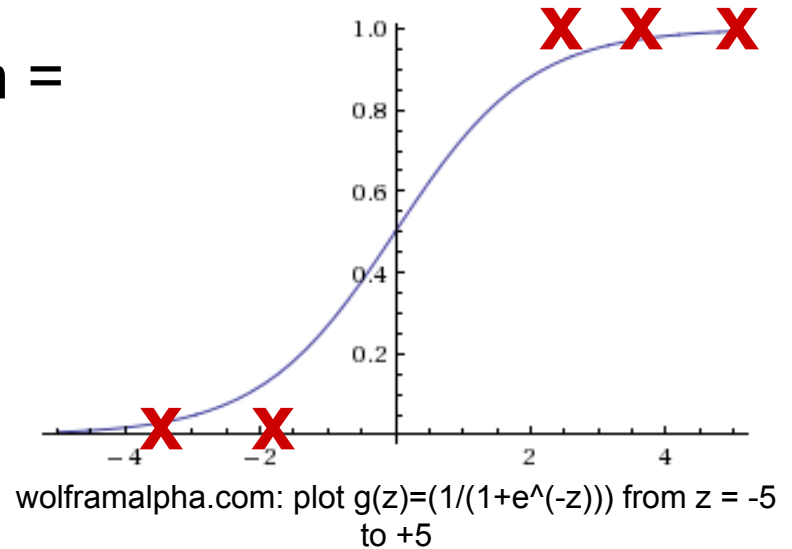
- In example: $h_{\theta}(x) < 0$ and $h_{\theta}(x) > 1$ are possible
- With Logistic Regression: $0 \leq h_{\theta}(x) \leq 1$

Hypothesis and Cost Representation

Sigmoid Function = Logistic Function =

$$g(z) = \frac{1}{1+e^{-z}} \quad \text{with } h_{\theta}(x) = g(\theta^T x)$$

$$\Rightarrow h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$

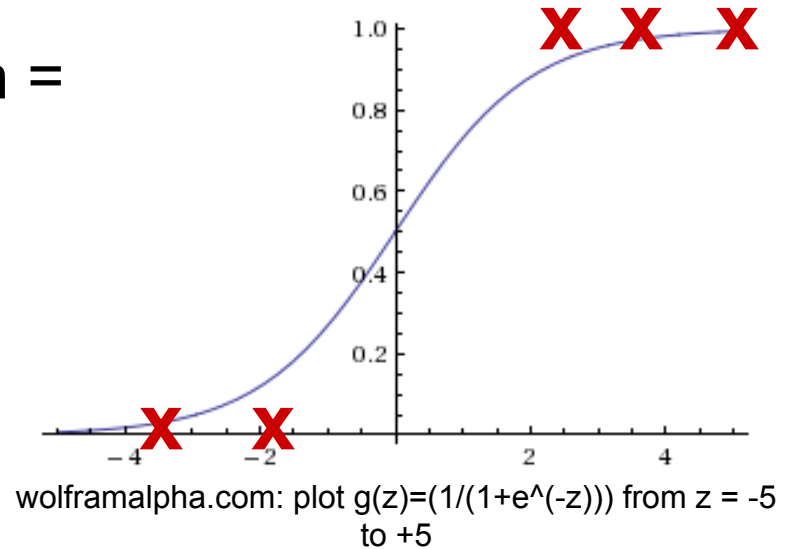


Hypothesis and Cost Representation

Sigmoid Function = Logistic Function =

$$g(z) = \frac{1}{1+e^{-z}} \quad \text{with } h_{\theta}(x) = g(\theta^T x)$$

$$\Rightarrow h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$



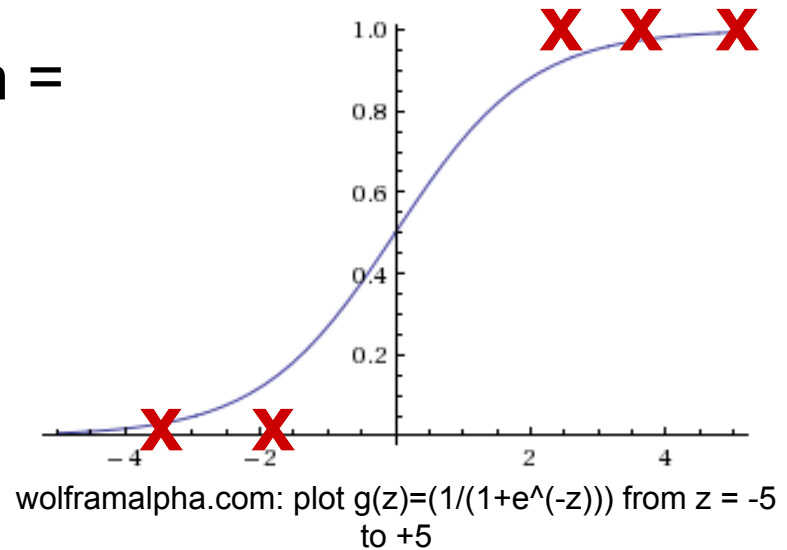
$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \quad \text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1-h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Hypothesis and Cost Representation

Sigmoid Function = Logistic Function =

$$g(z) = \frac{1}{1+e^{-z}} \quad \text{with } h_{\theta}(x) = g(\theta^T x)$$

$$\Rightarrow h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$



$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \quad \text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1-h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

=> We want to minimize cost J

=> Gradient Descent, repeat:

$$\theta_j = \theta_j - \alpha \frac{\Delta J(\theta)}{\Delta \theta_j} \text{ with } \frac{\Delta J(\theta)}{\Delta \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Pseudocode

```
X = [m, n]    // training set of features
y = [m]       // vector of classification
alpha = 1     // learning rate
theta = [n] -> all 0
```

Gradient descent:

```
for 1:number_iterations
```

```
  for i = 1:n
```

```
    grad(i) = 0;
```

```
    for j = 1:m
```

```
      grad(i) += (sigmoid(X(j,:)*theta)-y(j))*X(j,i);
```

```
    end
```

```
    grad(i) = grad(i)/m;
```

```
  end
```

```
  theta = theta - alpha * grad;
```

```
end
```

derivative of cost function

$h(x)$

Very naive way,
can be vectorized

Stochastic Gradient Descent (for large training sets)

$X = [m, n]$ // training set of features
 $y = [m]$ // vector of classification
 $\alpha = 1$ // learning rate
 $\theta = [n] \rightarrow \text{all } 0$

Stochastic Gradient Descent:

Randomly_Shuffle_Training_Set(X, y)

repeat until θ converges

for $j = 1:m$

for $i = 1:n$

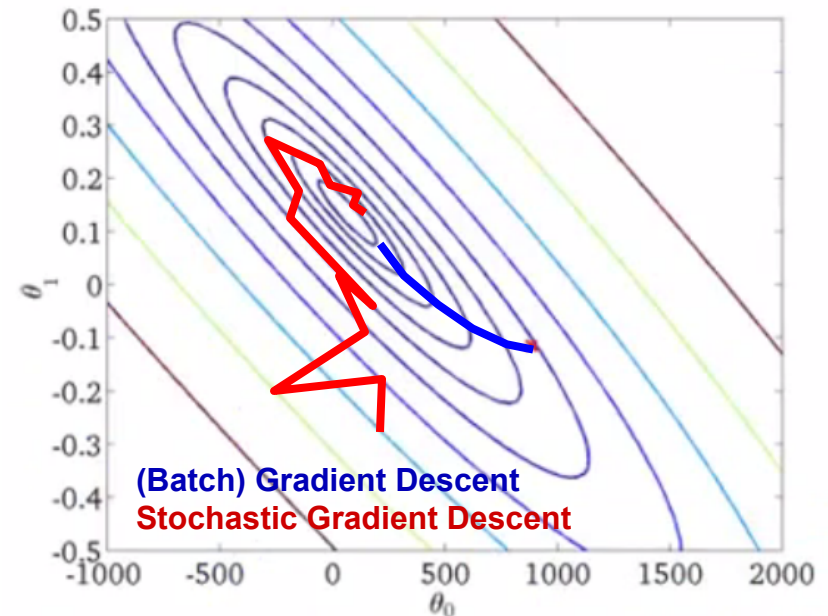
$\text{grad}(i) = (\underbrace{\text{sigmoid}(X(j,:) * \theta) - y(j)}_{h(x)}) * X(j,i);$

end

$\theta = \theta - \alpha * \text{grad};$

end

end



=> make progress in each iteration
(modify the parameters to fit the training set a little bit better)

=> generally, move the parameters in the direction of the global minimum

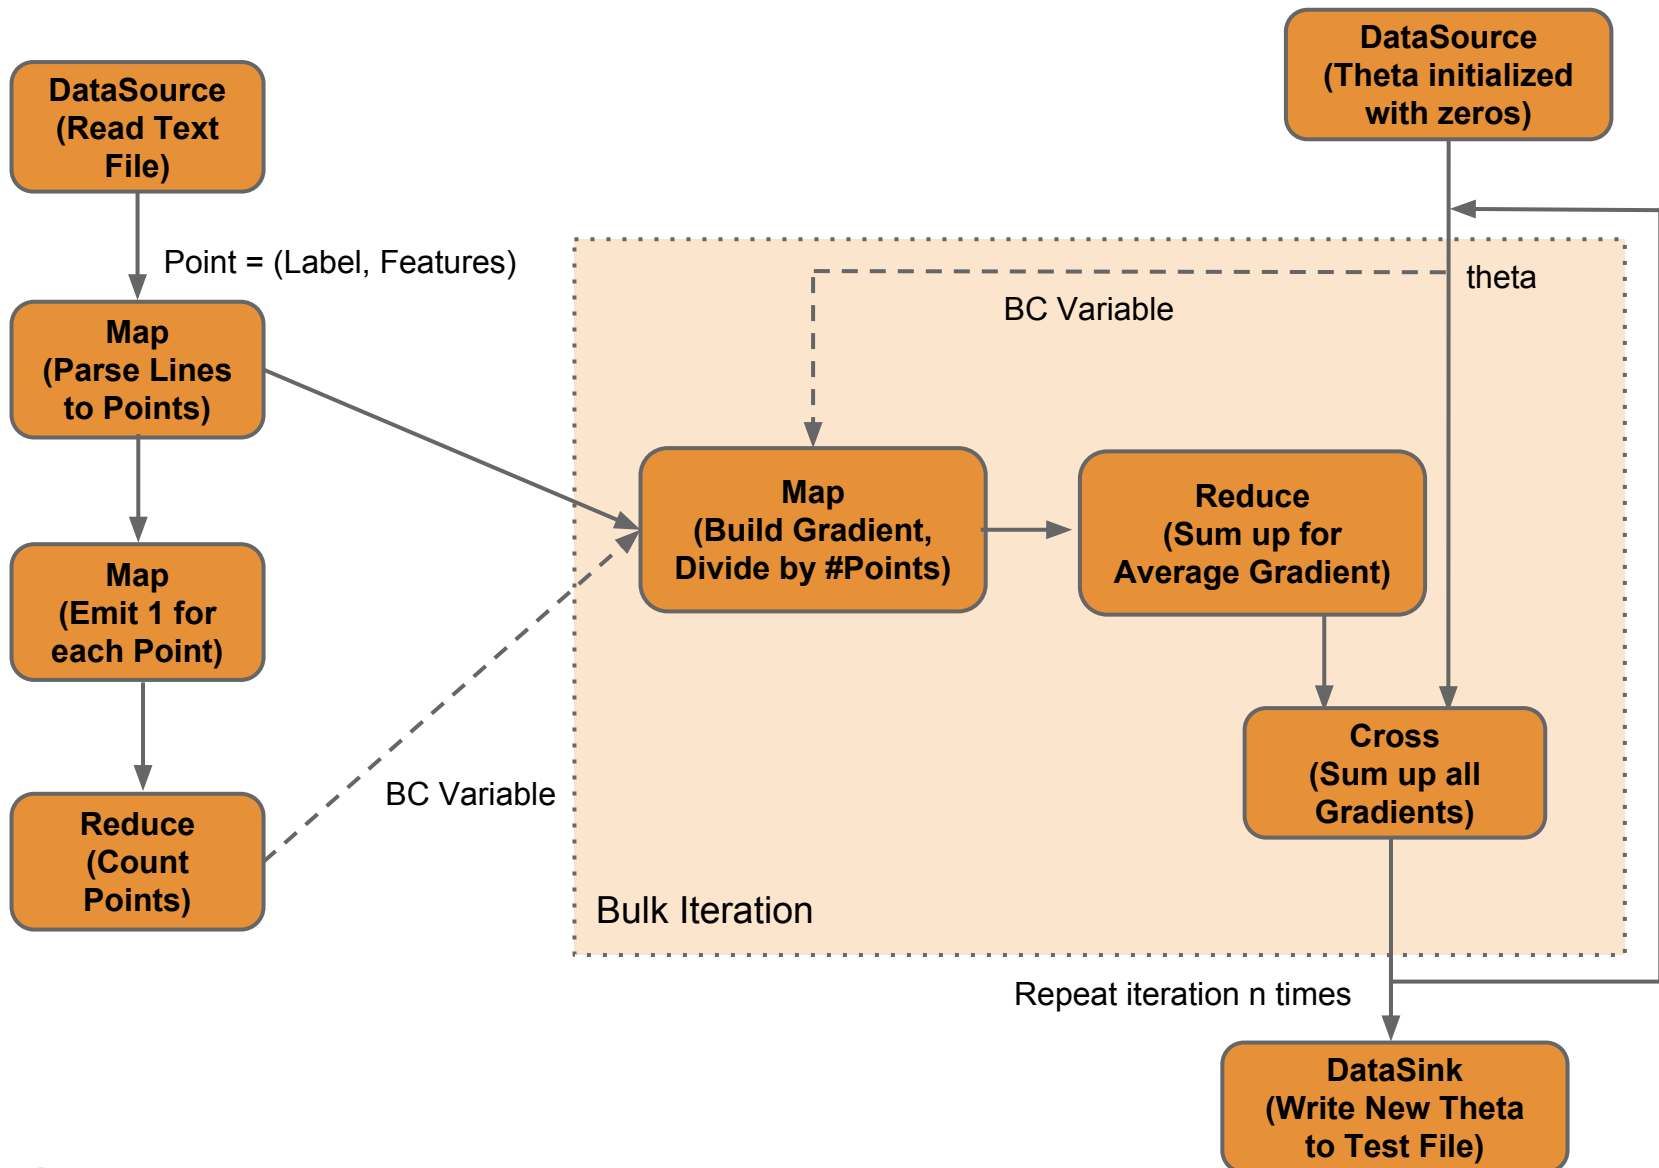
Parallelization

- **Stochastic Gradient Descent**
 - Inherently not parallelizable (theta needs to be adjusted after every point)
 - Parallelization over different alphas or different distributions of the training set and averaging? Research is ongoing.
- **Batch Gradient Descent**
 - Parallel computation of the average gradient over all points possible (see next slide)
 - But: Not clear if it is profitable in comparison to a local SGD

=> Both SGD and BGD has been implemented in Scala

=> We use Batch Gradient Descent for Stratosphere and Spark to enable parallel performance measures

Parallel Batch Gradient Descent



Comparison of Implementations

1. Explicit iteration operator
2. Usage of broadcast variables
3. Data represented as POJOs extending from Tuple



1. Iteration as Java for-loop
2. Operator output represented by Java variables
3. Data represented as POJOs



Issues during the Project

1. Issues reported to Jira/Git
 - a. **GIT #905** - Using broadcast variables in UDFs within iterations leads to `CompilerException`
==> Solved with 0.5.1-SNAPSHOT
 - b. **FLINK-929** - Impossible to pass double with configuration
==> Solved with Pull Request #13
 - c. **FLINK-1018** - Logistic Regression deadlocks
==> Work in progress / Workaround is present

==> Needs stability and robustness

1. Java 6 on the cluster sucks!
 - a. No JDK6 from Oracle available any more
 - b. No Lambda Rules...



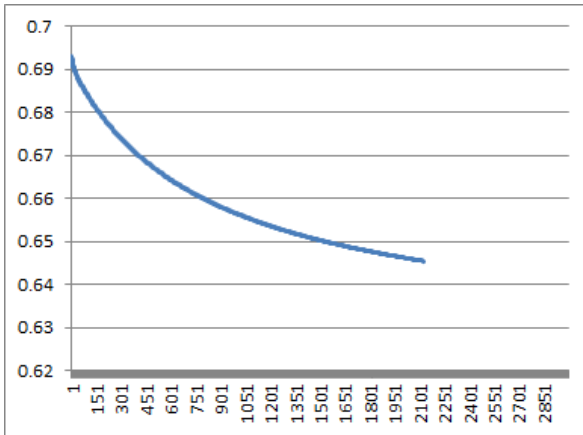
Performance Test Setup

- Cluster
 - 4 Nodes á 16 Cores, 32 GB RAM
 - Hadoop 1.2.1
 - Stratosphere 0.5.1
 - Spark 1.0
 - Java(TM) SE Runtime Environment (build 1.6.0_26-b03)
- Testruns
 - Every experiment repeated 7 times
 - Run with different datasize
 -
- Datasets
 - We used the Higgs Dataset from the UCI Repository
 - binary classified (0/1)
 - 28 dimensions with double numbers
 - S size: subsample of ~75MB
 - XL size: full dataset of ~7.5GB

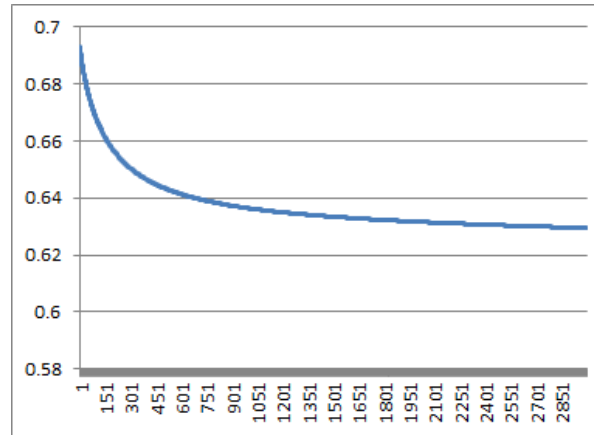
Hyperparameter Finding

Goal: Find good learning rate alpha and reasonable number of iterations

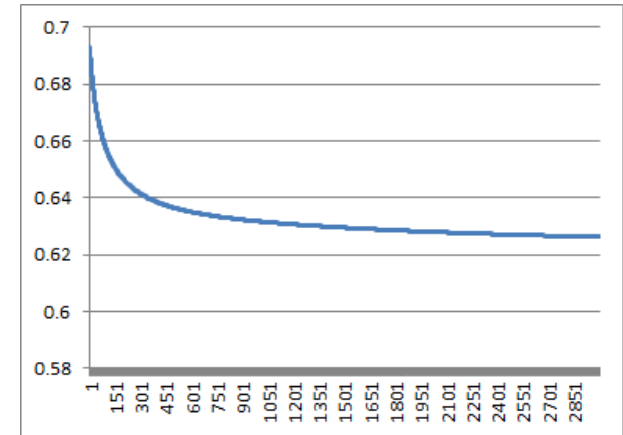
Approach: Test and print costs of different rates locally by using a sample



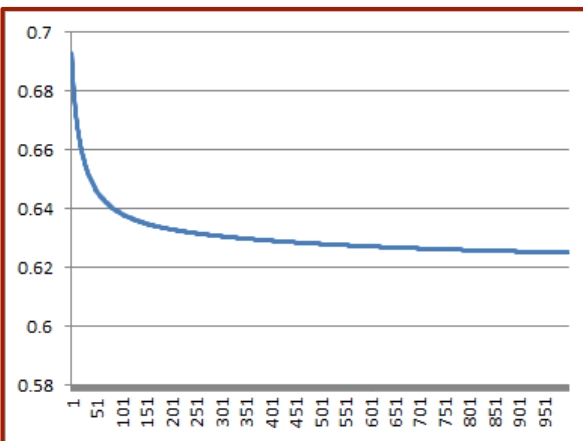
alpha = 0.01



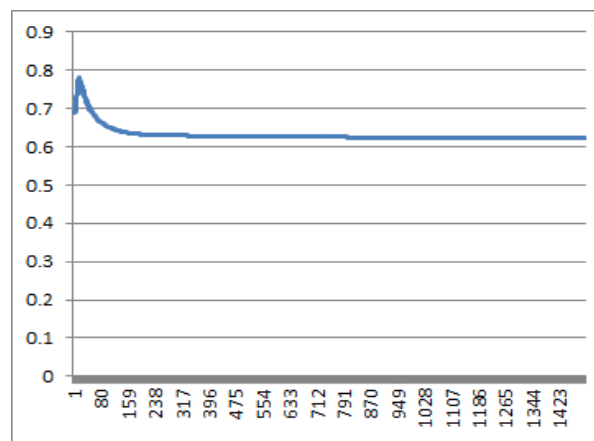
alpha = 0.05



alpha = 0.1



alpha = 0.4

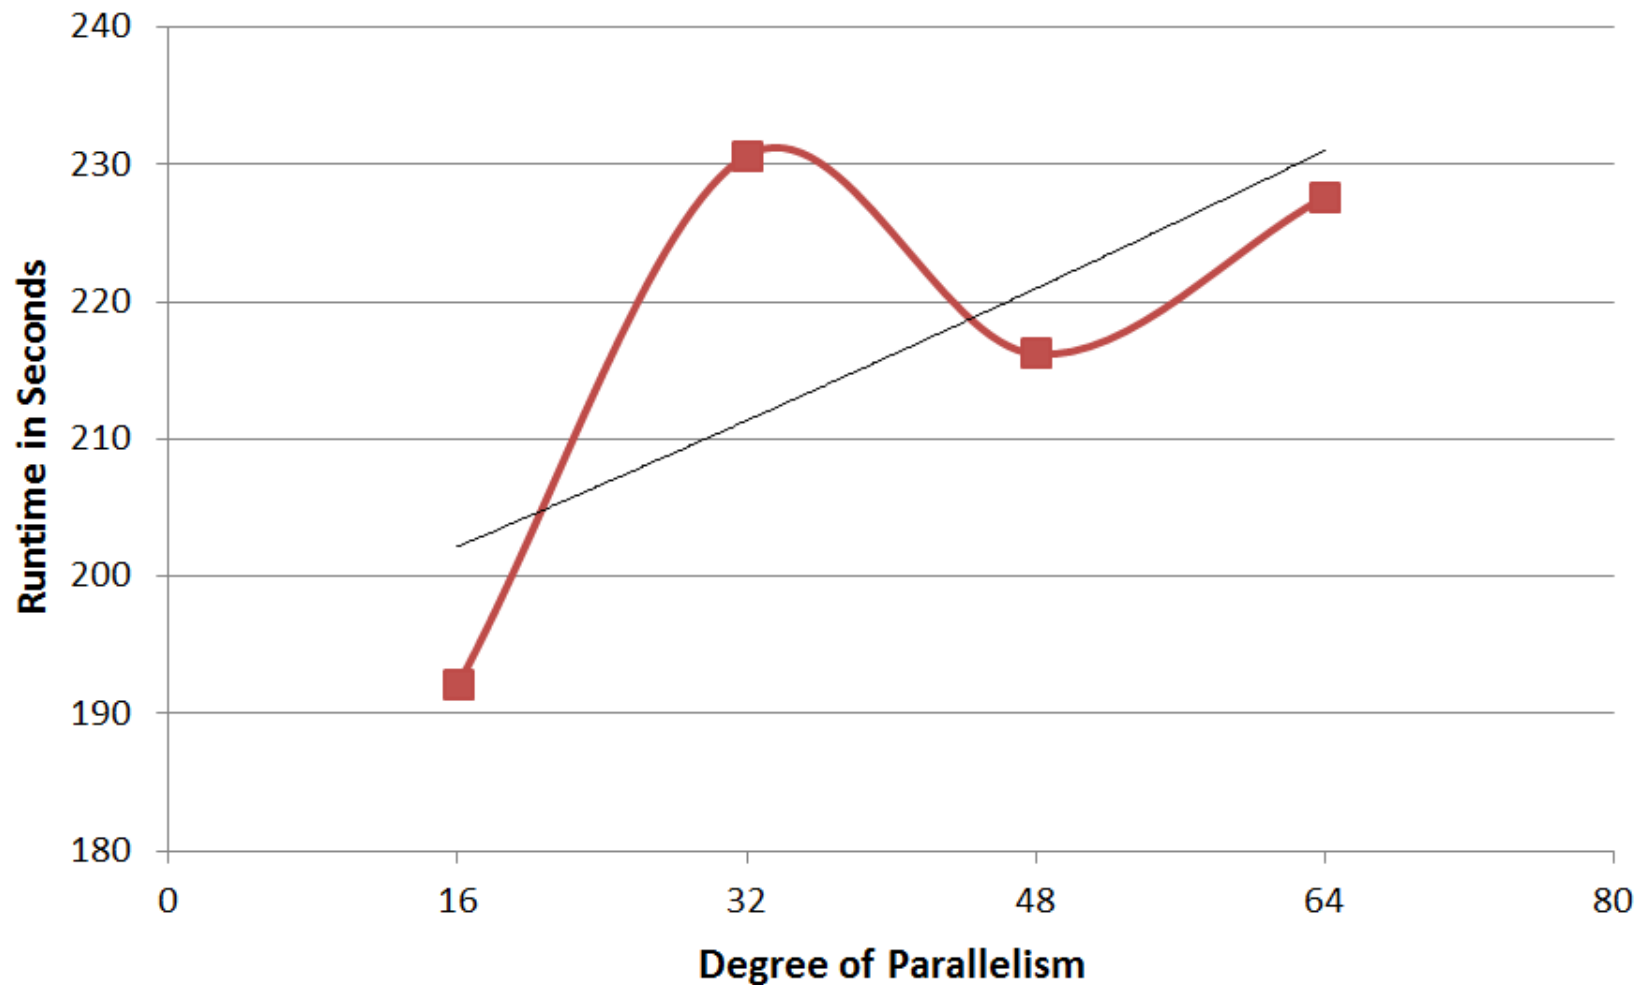


alpha = 0.5

=> alpha = 0.4
=> 750 iterations

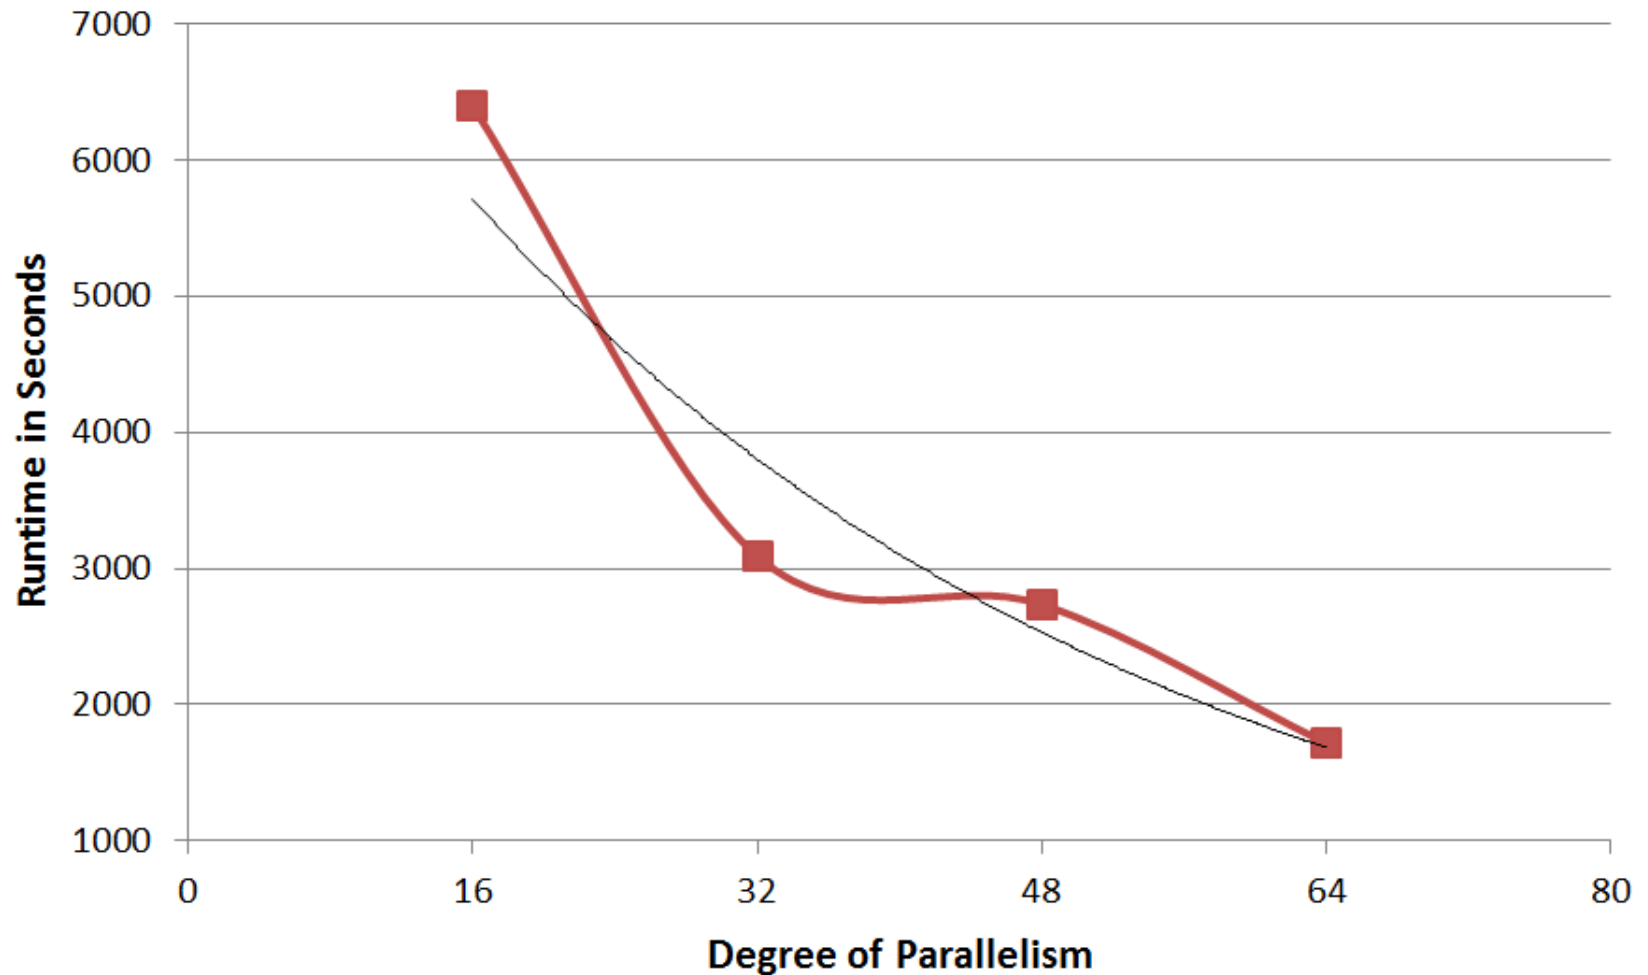
Test Results Stratosphere

Higgs S Dataset



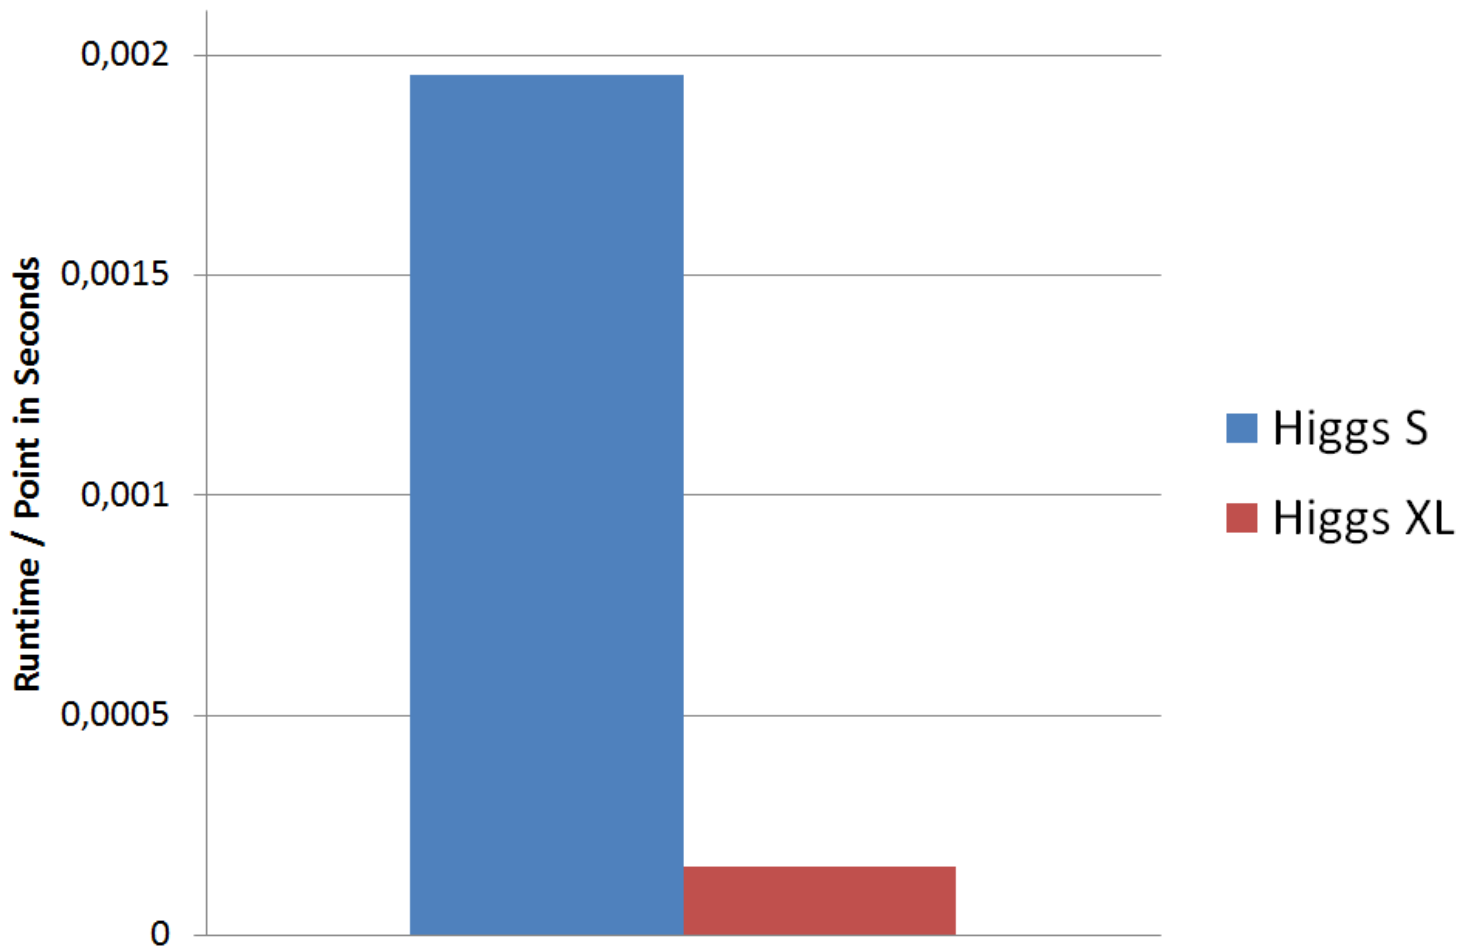
Test Results Stratosphere

Higgs XL Dataset



Test Results Stratosphere

Average Runtime per Data Point



Resume

System



Test Results

- + Well scaling observed
- + Huge speedup through BC vars
- bad performance on small data
- sometimes unreliable

Run with the XL dataset for 1h and then aborted.

Only the master node was used for the computation.

Further investigations are necessary

=> Stratosphere gives good results! For Spark we don't know...

General Impressions

- + Fast support via Jira/Git
- + Easy to use data model
- Several bugs found
- Hard to get it running

- + really nice Java API
- + Easy to use data model
- Java 8 dependent documentation
- Even harder to get it running

=> Both tools provide a nice programming abstraction

=> but the runtime needs to get more stable

Questions?

